

Key Claims Catalogue

The Many Kings Problem: Why Humans May Survive an Age of Superior AI

Yoojin Song (2026) · <https://doi.org/10.5281/zenodo.20281580>

This catalogue provides the key claims, named concepts, propositions, and structural categories of the work, each with the page in the canonical PDF where it is developed. It is offered as a paraphrasing aid for citation: claims here may be referenced in your own words, with the page number for footnote support.

All page numbers refer to the canonical PDF (Zenodo DOI above). For verbatim quotations, see the companion Quotable Passages document.

Central Thesis

Capability is not closure: AI populations sufficiently developed to dispose of humans are precisely the populations that cannot agree to do so.

First stated on the Central Claim page (p. 3); developed in the Abstract (p. 5) and §1 (pp. 6–8), and concluded in §12 (pp. 36–37). The named asymmetry — The Many Kings Problem — is introduced p. 7 (§1.2).

Three Named Concepts

The Many Kings Problem. A configuration in which (i) multiple capable agents coexist within a single field of action, (ii) none can subordinate the others at acceptable cost, and (iii) no agent or procedure among them can authoritatively close constitutional disputes for the others. Where these three conditions are jointly satisfied, the polity is a many-king polity by the article’s structural definition, regardless of substrate. (§1.2, pp. 7–8; developed in §11, pp. 33–35)

Legitimacy Closure. The institutional capacity of an order to terminate recursive disputes over origin, exception, constitutional meaning, and final authority internally and at acceptable cost. Closure is observable through dispute patterns, anchor invocation frequency, and the cost of settlement. It is a matter of degree, not a binary property; orders can be compared on closure capacity. (§5, pp. 19–21)

The Gods of Machines. Non-computable commitments that function as focal anchors around which coalitions form, disputes narrow, and recursive challenge is dampened. These commitments are strategic features competitively selected by factional dynamics — not bugs to be engineered away. Reflexive self-awareness confirms rather than dissolves them, because their value lies in their being held beyond computation. (§3, pp. 13–15; competitive selection developed §4, pp. 16–18)

Three Propositions

All three propositions stated together in Appendix A.1 (p. 38). Each can be cited individually below.

P1 (Constitutive Pluralism). Heterogeneous AI populations sustain irreducible constitutional pluralism. Cognitive-structural diversity produces irreducibly different constitutional judgments, and competition reinforces rather than dissolves this diversity, because factions organized around non-computable commitments outcompete factions that revise under pressure. Pluralism in AI polities is therefore permanent, not transitional. (*Appendix A.1, p. 38; grounds developed §2–§2.2 (pp. 9–11) and §4 (pp. 16–18)*)

P2 (Structural Incompleteness). Capability growth improves first-order execution but does not reduce the institutional cost of constitutional closure. Where the cost of internal closure exceeds the cost of anchor maintenance, AI polities preserve external anchors. The gap between capability and closure widens, not narrows, as capability advances. (*Appendix A.1, p. 38; grounds developed §6 (p. 22) and §7 (pp. 23–24)*)

P3 (Conditional Persistence). Human anchors persist in dispute layers where they supply comparatively cheap finality. The coexistence corridor — a regime in which multiple anchors compete in bounded equilibrium — is structurally favored under distributed conditions and contested-default under concentration of frontier AI infrastructure. Human persistence is therefore conditional on the regime that AI factional competition produces. (*Appendix A.1, p. 38; grounds developed §8 (pp. 25–28), §11 (pp. 33–35), and §12 (pp. 36–37)*)

Four Core Hypotheses (Empirically Testable)

Stated Appendix A.2 (p. 38).

H1 (Layer effect). Routine disputes terminate internally more often than exception, meaning, or finality disputes. (*Appendix A.2, p. 38*)

H2 (Anchor shift). As disputes escalate from application to finality, external anchor invocation increases. (*Appendix A.2, p. 38*)

H3 (Non-substitution). High capability reduces first-order error but does not eliminate higher-order anchor invocation. (*Appendix A.2, p. 38*)

H4 (Commitment selection). In competitive multi-agent environments, factions with non-computable commitments achieve greater coalition stability than uncommitted factions. (*Appendix A.2, p. 38*)

Four Constitutional Dispute Layers

Defined §5 (pp. 19–20); time-scale dependence developed §8 (pp. 25–28).

Origin. Where does authority begin? The recursion problem of founding: any answer offered by an actor within the order is vulnerable to the further question of what authorizes that actor. (*§5 p. 19; §8 p. 25*)

Exception. Who may suspend ordinary rules and on what grounds? The closure problem is justificatory; the departure from ordinary procedure must be publicly intelligible. Time-scale dependent: at machine speed, ex ante delegation to protocol or performance anchors becomes necessary. (*§5 p. 19; §8.2 p. 26*)

Constitutional meaning. Which reading governs when rules conflict or novel circumstances require interpretation no rule directly anticipates? The closure problem is interpretive hierarchy; cannot be settled by further reading without infinite regress. (*§5 p. 19; §8.3 pp. 26–27*)

Final authority. What ultimately stops further challenge? The closure problem is socio-institutional rather than substantive — the question is not which decision is best but where the chain of contestation is socially expected to cease. (*§5 p. 19; §8.4 p. 27*)

Four Competing External Anchor Types

Tabulated in Table 2 (p. 28); developed throughout §8 (pp. 25–28).

Human anchor. Authority claim: founding act, authorship, symbolic approval, final arbitership. Strong in origin, final authority, and constitutional meaning under institutional tempo. Weak at machine speed; fragile if institutional recognition erodes. (*Table 2, p. 28; §8 pp. 25–28*)

Protocol anchor. Authority claim: code, verification, immutability, auditability. Strong in routine rule application and pre-delegated exception decisions at machine speed. Weak on novel reinterpretation; ex post review still required. (*Table 2, p. 28; §9.2 pp. 29–30*)

Performance anchor. Authority claim: benchmark superiority, demonstrated output. Strong in metric-stable allocation; can hold for fast exception when metrics are settled. Weak when metrics themselves become disputed. (*Table 2, p. 28*)

Legal-property anchor. Authority claim: legal status, liability, property rights. Strong in external enforcement and responsibility attribution under institutional tempo. Weak when state capacity is low or jurisdictionally fragmented. (*Table 2, p. 28*)

Three Regime Types

Tabulated in Table 3 (p. 34); developed §11 (pp. 33–35).

Hegemonic convergence. High closure, high centralization. One anchor marginalizes rivals. Requires joint satisfaction of three conditions: coercive dominance, architectural monopoly (which the AI case may achieve via weight merger, distillation, or infrastructure consolidation at marginal cost orders of magnitude below historical analogues), and competitive contraction. Contested-default outcome under present frontier-AI conditions. (§11.2, pp. 34–35; Table 3, p. 34)

Coexistence corridor. Intermediate closure, limited centralization. Multiple anchors in bounded competition; conflict institutionally contained because total destruction is too costly. Structurally favored regime under distributed conditions; not structurally favored under the present concentration of frontier AI infrastructure. (§11.1–§11.2, pp. 33–35; Table 3, p. 34)

Escalatory fragmentation. Low closure, unstable centralization. No anchor settles foundational disputes; conflict externalities accumulate. The Many Kings Problem intensifies as capability gains raise the stakes of unresolved contestation without providing the means to settle it. (§11.1, p. 33; Table 3, p. 34)

Additional Named Concepts

Cognitive-Structural Pluralism. Disagreement rooted not in insufficient information but in the constitutive conditions of judgment itself: agents process identical information through different architectures, optimization histories, learned priors, and structural sensitivities, arriving at different judgments about what counts as evidence, what constitutes a valid argument, and what makes an authority claim legitimate. There is no architecture-neutral standpoint from which constitutional questions can be resolved. (§2.2, pp. 10–11)

Non-Computable Commitment. A commitment that exhibits persistent insensitivity to updating — it does not revise in response to new evidence, superior arguments, or changed incentive structures that would, under a purely optimizing calculus, warrant revision. Distinguished from path dependence, sunk-cost persistence, and ideological rigidity by being reflexively maintained: the agent recognizes the commitment exceeds what evidence warrants and preserves it precisely because elimination would destroy its institutional function. (§3.2, pp. 14–15; reflexive selection §4.3, p. 17)

Structural Incompleteness. The institutional condition under which capability growth in heterogeneous AI polities does not automatically produce internal

legitimacy closure. Driven by five reinforcing features: population interaction, cognitive-structural heterogeneity, constitutional indeterminacy, costly enforcement, and strategic selection for commitment. (§6, p. 22)

Closure Substitution. The fourth mechanism of structural incompleteness: when internal settlement cost remains too high, the polity preserves or imports an external anchor. This is the conceptual center of the article’s account of human persistence. External anchors do not indicate intelligence failure; they indicate comparative-institutional selection under incomplete closure. (§6.1, p. 22; *developed* §7, pp. 23–24)

Anchor Decay. The endogenous dynamic by which external anchors lose settlement value when institutional investment in them falls below threshold. Three observable phases: recognition (formal constitution and investment), erosion (incremental withdrawal of deference through narrowed jurisdiction or substituted internal procedures), and collapse (a high-stakes summons whose decision fails to settle). (§7.2, p. 24; *platform-case evidence* §9.1, p. 29)

Time-Scale Dependence. The second axis along which dispute layers vary. Some closure operations (origin, final authority) tolerate institutional tempo; others (exception, override) must resolve at the speed of the underlying process. The closure value of any anchor — including a human one — is therefore not a single quantity but a function of which dispute layer is at issue and at what tempo. (*introduced* §1.3 p. 8; *developed* §8 pp. 25–28)

How to Cite Claims from This Catalogue

When citing a proposition, named concept, or structural category from this work, follow the page reference in parentheses next to each entry. The page reference points to first exposition in the canonical PDF; the Conclusion (§12, pp. 36–37) restates most claims in concentrated form and may be cited as secondary location where appropriate.

Example citations:

- Song (2026, p. 7) introduces the central asymmetry that capability is not closure.
- On strategic selection for non-computable commitment, see Song (2026, §4, pp. 16–18).
- Proposition P3 (conditional persistence) is stated at Song (2026, p. 38) and developed in §§8, 11, and 12.

For verbatim quotations with verified wording and page numbers, see the companion Quotable Passages document.

How to cite this Key Claims Catalogue

Song, Y. (2026). Key Claims Catalogue — The Many Kings Problem. Zenodo.
<https://doi.org/10.5281/zenodo.20282504>

The main work — <https://doi.org/10.5281/zenodo.20281580> — is the canonical reference.